

▼ Predicting Antibiotic Resistance Genes | Proof Of Concept

The situation is a simple binary classification on sequence data. The Dataset was original gathered from kaggle, but the author privated the dataset soon after I downloaded (not kidding).

Results:

all models consistantly get above 90% accuracy with little optimization

- MLP = 94%
- GRU = 97%
- CNN = 98%

In summary, resistance is easily identified through genes, predicting based off Pan / meta genomes will be less simple. Many improvements could have been made in light of future models.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

import tensorflow as tf
from tensorflow import keras
from keras import layers

from functools import partial
from sklearn.metrics import ConfusionMatrixDisplay, accuracy_score
```

▼ Setup

```
from google.colab import drive
drive.mount('/content/drive')
!cd /content/drive/MyDrive/
!ls
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount()
drive sample_data



```
np_data = np.load('/content/drive/MyDrive/datasets/dataset.npy', allow_pickle=True)
np_data
```

```
array({'resistant': array([False, False, False, ..., True, True, False]), 'genes':
array(['ATGCACTACCGTATGATCCCCCTTCACTGGATGATGGAAATTGACTGCAATGGCTGCGCTAATAACTTTGTCTCGTC
```

```
'ATGCACTACCGTATGATTCACTGGATGATGGAAATTGACTGCAATGGCTGCGCTAATAATACTTTGTCTCGTCGTTGGAATTACGA
'ATGCACTACCGTATGATTCACTGGATCGTGATGGAAATTGACTGCAATGGCTGCGCTAATAATACTTTGTCTCGTCGTTGGAATTA
...,
'ATGCACTACCGTATGATTCACTGGATGATGGAACTGCAATGGCTGCGCTAATAATACTTTGTCTCGTCGTTGGAATTACGACTTC
'ATGCACTACCGTATGATTCACTGGATGATGGAAATTGACTTGCAATGGCTGCGCTAATAATACTTTGTCTCGTCGTTGGAATTACG
'ATGCACTACCGTATGATTCACTGGATGATGGAAATTGACTGCAATGGCTGCGCTAATAATACTTTGTCTCGTCGTTGGAATTACGA
dtype=object)}`, dtype=object)
```

```
# thanks to Alexander Scarlat .MD for the pandas dataframe conversion
Datadict = np_data[()]
df = pd.DataFrame.from_dict(Datadict)
print(df.shape)
df
```

```
(100000, 2)
```

| | resistant | gene |
|-------|-----------|---|
| 0 | False | ATGCACTACCGTATGATCCCCCTTCACTGGATGATGGAAATTGAA |
| 1 | False | ATGCACTACCGTATGATTCACTGGATGATGGAAATTGACTGCAA |
| 2 | False | ATGCACTACCGTATGATTCACTGGATCGTGATGGAAATTGACTG |
| 3 | False | ATGCACTACCGTATGATTCACTGGGTGATGGAAATTGACTGCAA |
| 4 | False | ATGCACGACCGTATCATTCACTGGATGATGGAAATTGACTGCAA |
| ... | ... | ... |
| 99995 | False | ATGCACTACCGTATGATTCACTGGATGATGGAACTCTCTTTGA |
| 99996 | False | ATGCACTACCGTATGATTCACTGGATGATGGAAATTGACTGCAA |
| 99997 | True | ATGCACTACCGTATGATTCACTGGATGATGGAACTGCAATGGCT |
| 99998 | True | ATGCACTACCGTATGATTCACTGGATGATGGAAATTGACTTGCA |

```
# gene
X = df['genes'].iloc[0]
print(len(X), X)
```


```
156 ATGCACTACCGTATGATCCCCCTTCACTGGATGATGGAAATTGACTGCAATGGCTGCGCTAATAATACTTTGTCTCGTCGTTG
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 2 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   resistant  100000 non-null  bool
 1   genes      100000 non-null  object
dtypes: bool(1), object(1)
memory usage: 879.0+ KB
```

```
val_counts = df["resistant"].value_counts()
print(val_counts)
val_counts.plot(kind="bar")
```

```
False    50297
True      49703
Name: resistant, dtype: int64
<Axes: >
```



```
from keras.utils import pad_sequences

vocab = ['A', 'C', 'G', 'T']
char2idx = {
    char: (idx+1) for idx, char in enumerate(vocab)
}
```

```

# tokenize (convert to ATGC to 1-4)
def tokenize(sequence):
    return [char2idx[char] for char in sequence]

# one hot encode (convert to sparse vector)
def encode(tokens):
    tensor = tf.constant(tokens)
    one_hot = np.asarray(tf.one_hot(tensor, len(vocab)))
    return one_hot.astype("int32")

df['genes_tokenized'] = df['genes'].apply(tokenize)

# pad right with 0's for the tokenized column
X_token = pad_sequences(df['genes_tokenized'], maxlen=160, padding='post', truncating='post',

X_enc = np.array([encode(tokens) for tokens in X_token])

print(df.head(3))
print("OHE:\n ", X_enc[0][0][:5], "\n .... ")

```

| | resistant | genes \ |
|---|-----------|--|
| 0 | False | ATGCACTACCGTATGATCCCCCTTCACTGGATGATGGAAATTGACT... |
| 1 | False | ATGCACTACCGTATGATTCACCTGGATGATGGAAATTGACTGCAATG... |
| 2 | False | ATGCACTACCGTATGATTCACCTGGATCGTGATGGAAATTGACTGCA... |

| | genes_tokenized |
|---|--|
| 0 | [1, 4, 3, 2, 1, 2, 4, 1, 2, 2, 3, 4, 1, 4, 3, ...] |
| 1 | [1, 4, 3, 2, 1, 2, 4, 1, 2, 2, 3, 4, 1, 4, 3, ...] |
| 2 | [1, 4, 3, 2, 1, 2, 4, 1, 2, 2, 3, 4, 1, 4, 3, ...] |

OHE:

| |
|-----------|
| [0 1 0 0] |
| |

```

y = np.array(df['resistant']).astype('int32')

def split(X, ratio=0.8):
    # take out test set (top 20%)
    test_split = int(len(X) * ratio)
    full_train, test = X[:test_split], X[test_split:]

    # hold out top 20% of training for per epoch validation
    validation_split = int(len(full_train) * ratio)
    train, validation = full_train[:validation_split], full_train[validation_split:]

    return train, validation, test

y_train, y_valid, y_test = split(y)
X_enc_train, X_enc_valid, X_enc_test = split(X_enc)
X_token_train, X_token_valid, X_token_test = split(X_token)

```

▼ SVM

```
''' from sklearn.metrics import ConfusionMatrixDisplay, accuracy_score
from sklearn.svm import SVC

svm = SVC(kernel='linear')
svm.fit(X_token_train, y_train)

y_pred = svm.predict(X_token_test)

accuracy = accuracy_score(y_test, y_pred)
print("Accuracy: ", accuracy)

ConfusionMatrixDisplay.from_predictions(y_test, y_pred) '''

# colab dies -> to much data for non parametric

' from sklearn.metrics import ConfusionMatrixDisplay, accuracy_score\nfro
m sklearn.svm import SVC\n\nsvm = SVC(kernel='linear')\nsvm.fit(X_token
train, y_train)\n\ny_pred = svm.predict(X_token_test)\n\naccuracy = accu
```

▼ MLP

```
DefaultDense = partial(
    layers.Dense,
    activation='relu',
    kernel_initializer="he_uniform",
    kernel_regularizer= keras.regularizers.l2(0.01)
)
mlp_model = keras.Sequential([
    keras.Input(160),
    DefaultDense(20), # technically not a MLP cos 2 hidden layers but who cares
    DefaultDense(10),
    layers.Dense(1, activation='sigmoid')
])
mlp_model.compile(
    loss='binary_crossentropy',
    optimizer='adam',
    metrics=['accuracy']
)
mlp_model.summary()
```

Model: "sequential_5"

| Layer (type) | Output Shape | Param # |
|-----------------|--------------|---------|
| dense_5 (Dense) | (None, 20) | 3220 |

| | | |
|-----------------|------------|-----|
| dense_6 (Dense) | (None, 10) | 210 |
| dense_7 (Dense) | (None, 1) | 11 |

```
=====
Total params: 3,441
Trainable params: 3,441
Non-trainable params: 0
=====
```

```
early_stopping_cb = keras.callbacks.EarlyStopping(patience=3) # stop training if overfits
mlp_history = mlp_model.fit(X_token_train, y_train, epochs=7, batch_size=32,
                           validation_data=(X_token_valid, y_valid),
                           callbacks=[early_stopping_cb])
```

```
Epoch 1/7
2000/2000 [=====] - 11s 4ms/step - loss: 0.5642 - accuracy: 0.
Epoch 2/7
2000/2000 [=====] - 7s 4ms/step - loss: 0.3807 - accuracy: 0.8
Epoch 3/7
2000/2000 [=====] - 7s 4ms/step - loss: 0.3240 - accuracy: 0.8
Epoch 4/7
2000/2000 [=====] - 8s 4ms/step - loss: 0.2894 - accuracy: 0.9
Epoch 5/7
2000/2000 [=====] - 7s 4ms/step - loss: 0.2677 - accuracy: 0.9
Epoch 6/7
2000/2000 [=====] - 8s 4ms/step - loss: 0.2536 - accuracy: 0.9
Epoch 7/7
2000/2000 [=====] - 7s 3ms/step - loss: 0.2382 - accuracy: 0.9
```



```
pd.DataFrame(mlp_history.history)[['accuracy', 'val_accuracy']].plot()
```

<Axes: >



```

y_pred = mlp_model.predict(X_token_test)
y_pred = (y_pred > 0.5).astype("int32")
mlp_accuracy = accuracy_score(y_test, y_pred)

```

```

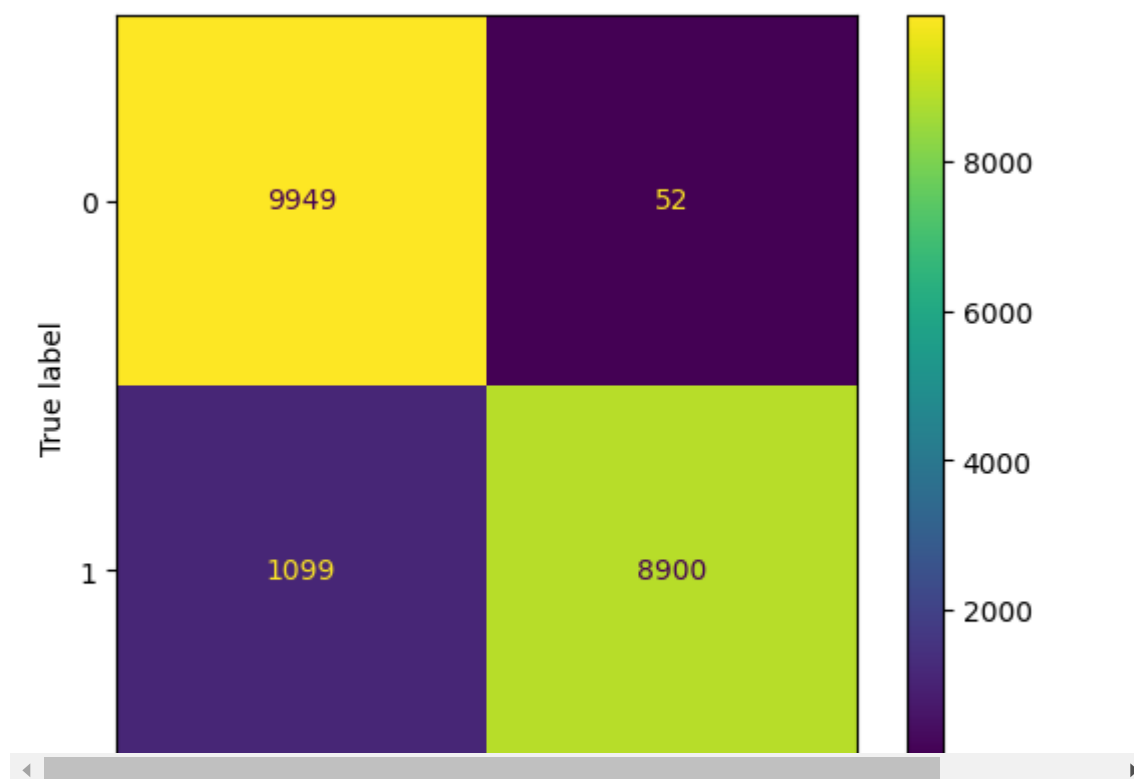
print("Accuracy: ", mlp_accuracy)
print(ConfusionMatrixDisplay.from_predictions(y_test, y_pred))

```

625/625 [=====] - 3s 4ms/step

Accuracy: 0.94245

<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x...



GRU

```

gru_model = keras.Sequential([
    layers.Bidirectional(
        # bidirection wrapper converts LSTM to GRU
        layers.LSTM(100, kernel_initializer='he_normal',
                    # kernel_regularizer='l2',
                    # dropout=0.3,
                    activation='tanh'),
    )
])

```

```

        input_shape=(160, 4)
    ),
    # add another dense?
    layers.Dense(1, activation='sigmoid')
])
gru_model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
gru_model.summary()

```

Model: "sequential_4"

| Layer (type) | Output Shape | Param # |
|---------------------------------|--------------|---------|
| bidirectional_5 (Bidirectional) | (None, 200) | 84000 |
| dense_4 (Dense) | (None, 1) | 201 |

=====
 Total params: 84,201
 Trainable params: 84,201
 Non-trainable params: 0
 =====

```

early_stopping_cb = keras.callbacks.EarlyStopping(patience=3)
gru_history = gru_model.fit(X_enc_train, y_train, epochs=7, batch_size=32,
                           validation_data=(X_enc_valid, y_valid),
                           callbacks=[early_stopping_cb])

```

```

Epoch 1/7
2000/2000 [=====] - 39s 18ms/step - loss: 0.4432 - accuracy: 0
Epoch 2/7
2000/2000 [=====] - 36s 18ms/step - loss: 0.2597 - accuracy: 0
Epoch 3/7
2000/2000 [=====] - 35s 17ms/step - loss: 0.2255 - accuracy: 0
Epoch 4/7
2000/2000 [=====] - 34s 17ms/step - loss: 0.1987 - accuracy: 0
Epoch 5/7
2000/2000 [=====] - 38s 19ms/step - loss: 0.3001 - accuracy: 0
Epoch 6/7
2000/2000 [=====] - 31s 15ms/step - loss: 0.1954 - accuracy: 0
Epoch 7/7
2000/2000 [=====] - 30s 15ms/step - loss: 0.1679 - accuracy: 0

```




```

pd.DataFrame(gru_history.history)[['accuracy', 'val_accuracy']].plot()

```


<Axes: >



overfits; should re-add dropout/regularization

```
y_pred = gru_model.predict(X_enc_test)
y_pred = (y_pred > 0.5).astype("int32")
gru_accuracy = accuracy_score(y_test, y_pred)

print("Accuracy: ", gru_accuracy)
print(ConfusionMatrixDisplay.from_predictions(y_test, y_pred))
```

625/625 [=====] - 5s 6ms/step
Accuracy: 0.9799

▼ CNN

```
DefaultConv = partial(layers.Conv1D, kernel_size=3, strides=1, use_bias=False,
                      padding='same', kernel_initializer='he_normal')

cnn_model = keras.Sequential([

    DefaultConv(32, strides=1, input_shape=(160, 4)),

    layers.Dropout(0.4),

    layers.MaxPooling1D(3, strides=2),

    DefaultConv(64),

    layers.MaxPooling1D(2, strides=2),

    layers.Flatten(),

    layers.Dropout(0.4),

    layers.Dense(64, activation="relu"),

    layers.Dense(1, activation="sigmoid")

])

cnn_model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
cnn_model.summary()
```

Model: "sequential_7"

| Layer (type) | Output Shape | Param # |
|--------------------------------|-----------------|---------|
| ===== | | |
| conv1d_6 (Conv1D) | (None, 160, 32) | 384 |
| dropout_6 (Dropout) | (None, 160, 32) | 0 |
| max_pooling1d_2 (MaxPooling1D) | (None, 79, 32) | 0 |
| conv1d_7 (Conv1D) | (None, 79, 64) | 6144 |
| max_pooling1d_3 (MaxPooling1D) | (None, 39, 64) | 0 |
| flatten_4 (Flatten) | (None, 2496) | 0 |

| | | |
|---------------------|--------------|--------|
| dropout_7 (Dropout) | (None, 2496) | 0 |
| dense_11 (Dense) | (None, 64) | 159808 |
| dense_12 (Dense) | (None, 1) | 65 |

```
=====
Total params: 166,401
Trainable params: 166,401
Non-trainable params: 0
```


```
early_stopping_cb = keras.callbacks.EarlyStopping(patience=3)
cnn_history = cnn_model.fit(X_enc_train, y_train, epochs=15, batch_size=32,
                           validation_data=(X_enc_valid, y_valid),
                           callbacks=[early_stopping_cb])
```

```
Epoch 1/15
2000/2000 [=====] - 10s 4ms/step - loss: 0.3479 - accuracy: 0.
Epoch 2/15
2000/2000 [=====] - 9s 4ms/step - loss: 0.2311 - accuracy: 0.8
Epoch 3/15
2000/2000 [=====] - 8s 4ms/step - loss: 0.1686 - accuracy: 0.9
Epoch 4/15
2000/2000 [=====] - 8s 4ms/step - loss: 0.1372 - accuracy: 0.9
Epoch 5/15
2000/2000 [=====] - 9s 5ms/step - loss: 0.1126 - accuracy: 0.9
Epoch 6/15
2000/2000 [=====] - 8s 4ms/step - loss: 0.0960 - accuracy: 0.9
Epoch 7/15
2000/2000 [=====] - 9s 4ms/step - loss: 0.0875 - accuracy: 0.9
Epoch 8/15
2000/2000 [=====] - 8s 4ms/step - loss: 0.0834 - accuracy: 0.9
Epoch 9/15
2000/2000 [=====] - 8s 4ms/step - loss: 0.0774 - accuracy: 0.9
Epoch 10/15
2000/2000 [=====] - 9s 4ms/step - loss: 0.0732 - accuracy: 0.9
```



```
pd.DataFrame(cnn_history.history)[['accuracy', 'val_accuracy']].plot()
```

<Axes: >



```

y_pred = cnn_model.predict(X_enc_test)
y_pred = (y_pred > 0.5).astype("int32")
cnn_accuracy = accuracy_score(y_test, y_pred)


print("Accuracy: ", cnn_accuracy)
print(ConfusionMatrixDisplay.from_predictions(y_test, y_pred))

```

625/625 [=====] - 1s 2ms/step

Accuracy: 0.97245

<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x7f7a9c16f340



Improvements that could have been made

- padding sequences is lazy and not suitable for GRU; real scenario would make use of variable length sequences
- Hyper params werent optimized
- didnt use codons

